

# Data Plane Programmability

## Introduction to P4



Tomasz Osiński, PhD



# Agenda

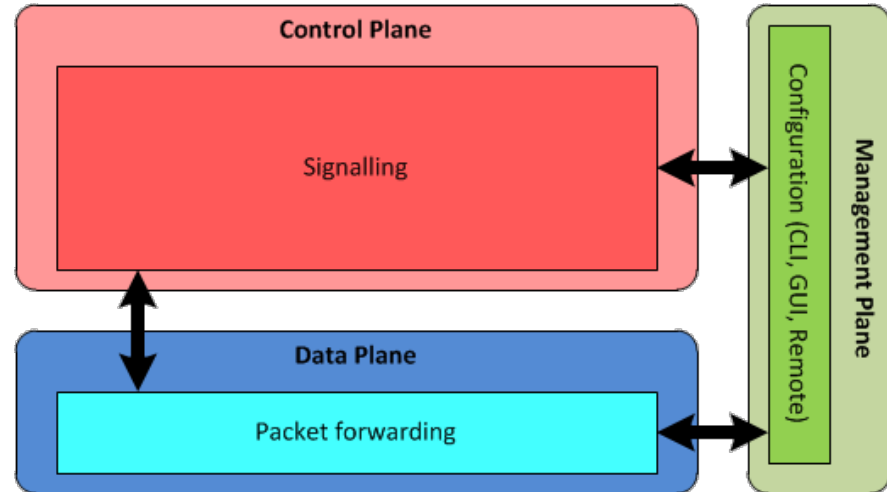
- **Introduction to Data Plane Programmability**
  - **Motivation**
  - **Concept**
  - **Relationship with OpenFlow & SDN**
- **P4**
  - **Architecture**
  - **Language**
  - **Toolchain**
  - **Demo**
- **P4Runtime protocol**
- **Use cases**
- **Research activities**
- **Summary**

# Basics

# Telecommunications Architecture (quick reminder)

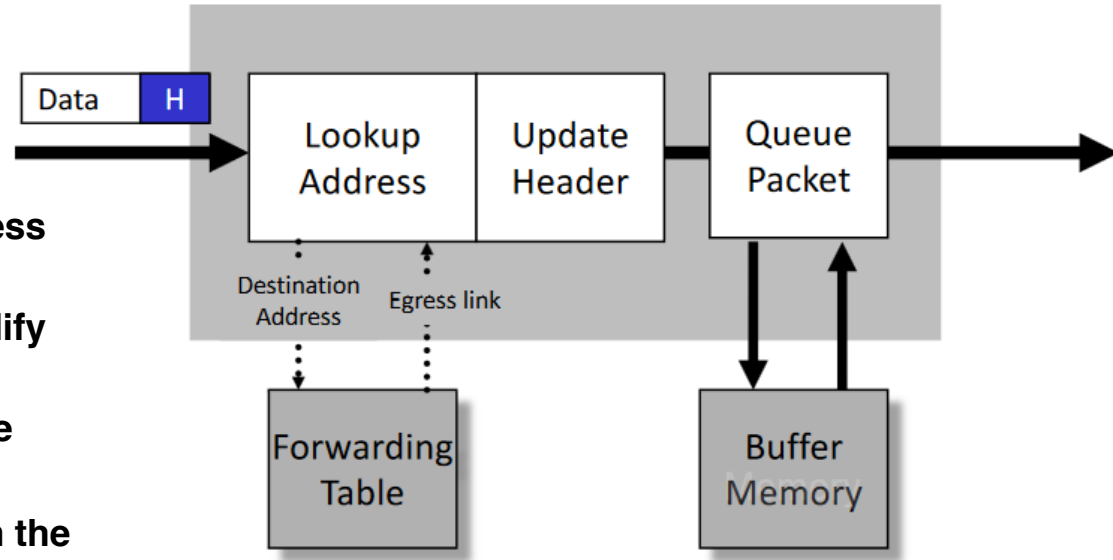
- **Classical architecture consists of 3 planes**
  - Group of protocols
- **Differences between planes:**
  - Kind of traffic
  - Performance
  - Programming languages
  - Different execution environments

Exercise!



# What the Data Plane is?

- The mechanism responsible for packet processing based on the protocol's headers
- Operations:
  1. **Parse** – read the packet
  2. **Lookup** – determine how to process the packet
  3. **Update Header** (optionally) – modify packet header's field(s)
  4. **Switch/Route** – send packet to the output port
  5. **Queue Packet** – put the packet on the output queue
  6. **Deparse** – create packet and send on the wire



# Internet Router – the real-world example

## 1. Parse

- Examine IP packet – read the packet's header and determine packet length

## 2. Lookup

- Check how to forward packet based on destination IP address

## 3. Update Header

- Decrement TTL, calculate the IP checksum

## 4. Switch/Route packet

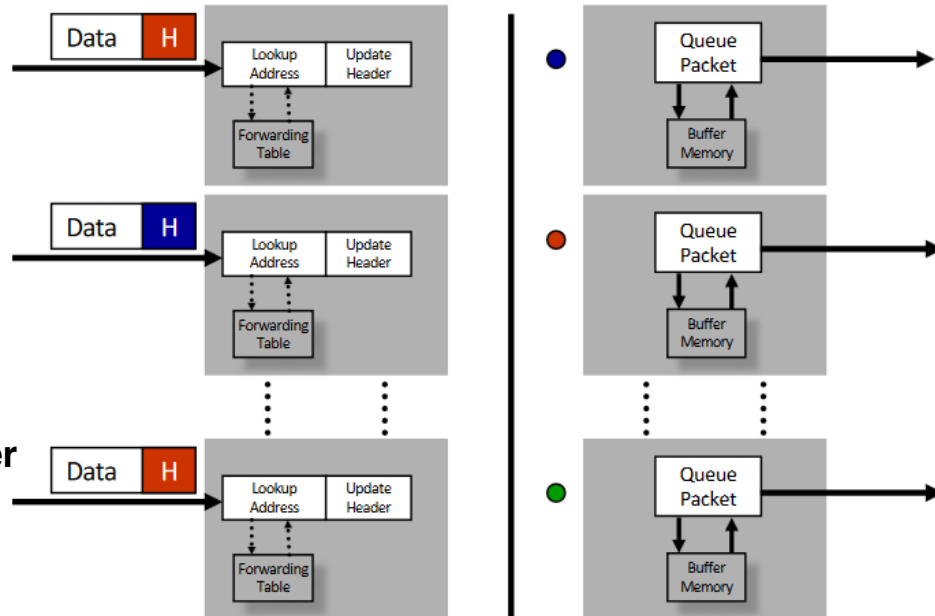
- Forward to the correct egress port

## 5. Queue packet

- Find the Ethernet DA for the next hop router
- Put the packet on the egress port's queue

## 6. Deparse


- Create Ethernet frame and send it



# Introduction to Data Plane Programmability

# Software-Defined Networking

- **Software-Defined Networking (SDN)**
  - Data plane and control plane decoupling
  - Logically centralized control plane
  - Programmable control plane and Open API to data plane
- **OpenFlow protocol**
- **Issues:**
  - Data plane protocol evolution requires changes to standards (e.g. OpenFlow specification)
  - Data plane scalability (due to TCAM)

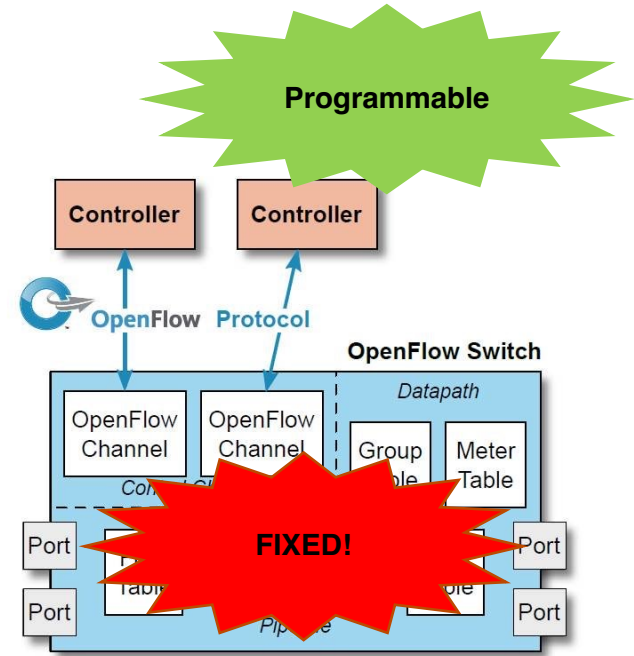




# Why Data Plane Programmability?

- **Limitations of OpenFlow:**
  - **Examples:**
    - Support for GTP protocol (vEPC)
    - Tunnelling protocols (e.g. MPLS over UDP)
    - Service Function Chaining (NSH headers)
- **Vendor lock-in (oops?)**

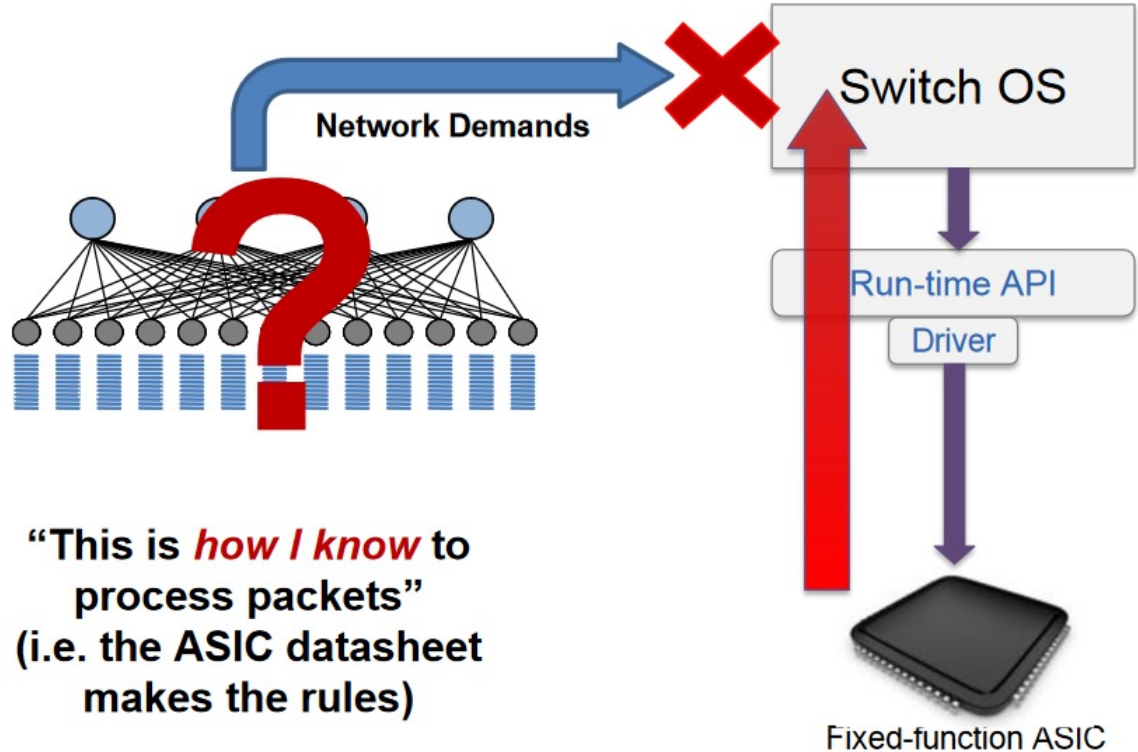
Version	Date	Header Fields
OF 1.0	Dec 2009	12 fields (Ethernet, TCP/IPv4)
OF 1.1	Feb 2011	15 fields (MPLS, inter-table metadata)
OF 1.2	Dec 2011	36 fields (ARP, ICMP, IPv6, etc.)
OF 1.3	Jun 2012	40 fields
OF 1.4	Oct 2013	41 fields



# Motivation (1/2)

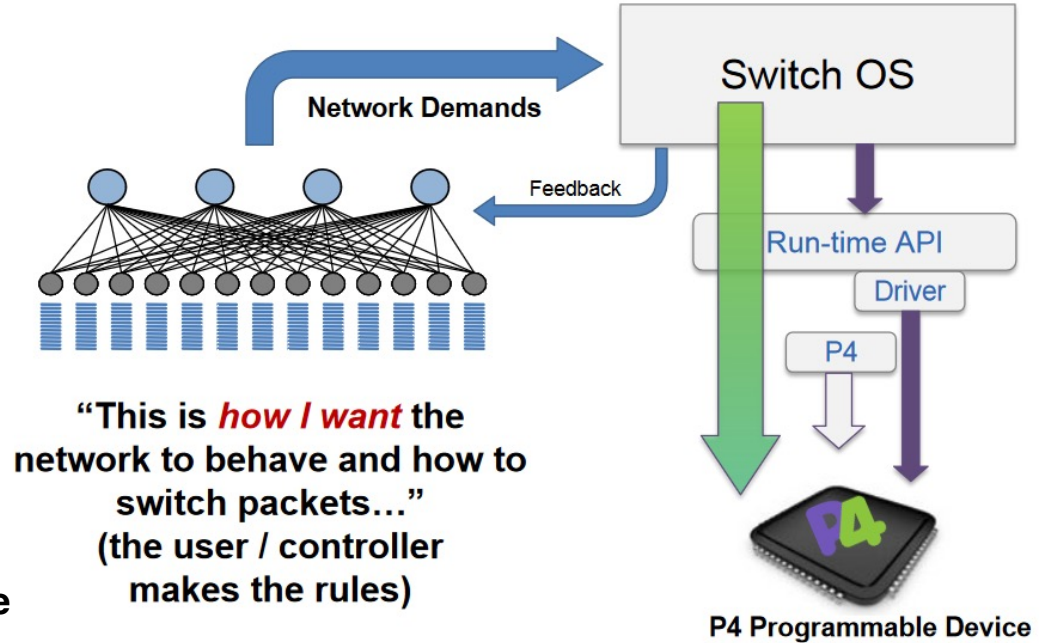
## Status quo: Bottom-Up design

- Slow release cycle (for ASICs ~4 years)
  - VxLAN case study
    - Standard: 2010 by Cisco & VMWare
    - Implementation: 2014



## Motivation (2/2)

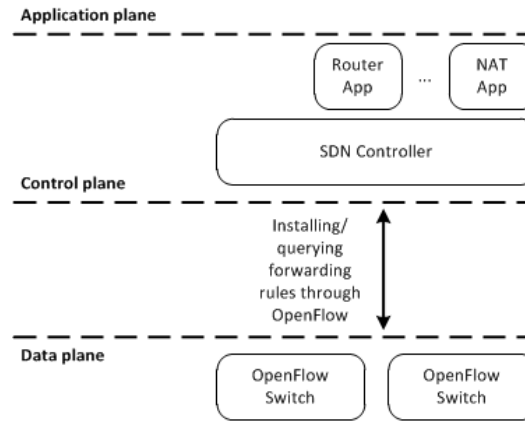
### A better approach: Top-Down approach



Why Does the Internet Need a Programmable Forwarding Plane by Nick McKeown  
<https://www.youtube.com/watch?v=zR88Nlg3n3g>

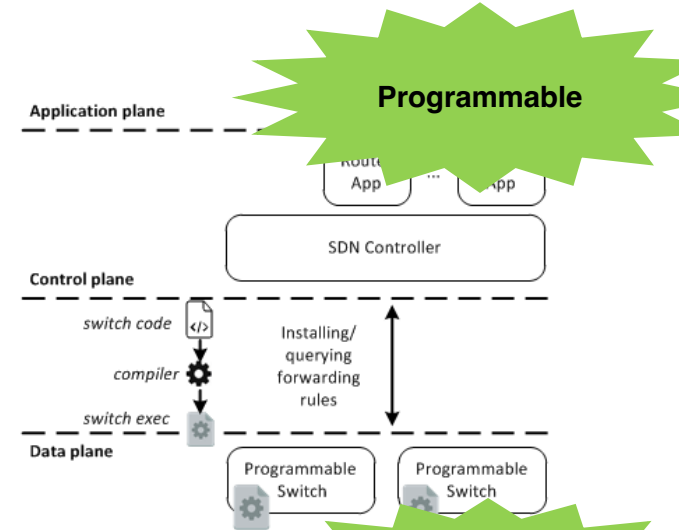
# What is Data Plane Programmability?

- Allows to define data plane of device in software
- Fixed OpenFlow switch → Programmable Switch
- Open API to program data plane (change software version of data plane device)



a)

Classical OpenFlow



b)

Programmable data plane

# Benefits of Data Plane Programmability

- **New Features** – Fast Time to Market (TTM) for new protocols
- **Software development style:**
  - **Debugging**
  - **Bug fixing**
  - **Software re-use (libraries)**
  - **Software upgrade**
  - **CI/CD**
- **Reduced complexity** – Remove unused protocols
- **Increased reliability and security** – reduced risk by removing unused protocols
- **New Use Cases** – e.g. telemetry
- **You keep your own ideas..** – Intellectual Property (IP)

# **P4 – Programming Protocol-independent Packet Processors**

P4 is a high-level programming language for Software-defined Networking (SDN).

It is intended to describe the behavior of the data plane of devices that forwards, modifies or inspects network traffic.



## P4 – a brief history

- **July 2014** – First paper „P4 – Programming Protocol-Independent Packet Processors”
- **September 2014** – First P4<sub>14</sub> specification
- **May 2017** – P4<sub>16</sub> specification

<http://github.com/p4lang/>

### P4: Programming Protocol-Independent Packet Processors

Pat Bosshart<sup>†</sup>, Dan Daly<sup>\*</sup>, Glen Gibb<sup>†</sup>, Martin Izzard<sup>†</sup>, Nick McKeown<sup>‡</sup>, Jennifer Rexford<sup>\*\*</sup>,  
Cole Schlesinger<sup>\*\*</sup>, Dan Talayco<sup>†</sup>, Amin Vahdat<sup>¶</sup>, George Varghese<sup>§</sup>, David Walker<sup>\*\*</sup>  
<sup>†</sup>Barefoot Networks   <sup>\*</sup>Intel   <sup>‡</sup>Stanford University   <sup>\*\*</sup>Princeton University   <sup>¶</sup>Google   <sup>§</sup>Microsoft Research



# The P4 organization



## P4.org Membership



Original P4 Paper Authors:



Operators/  
End Users



Systems



Targets



Solutions/  
Services



Academia/  
Research



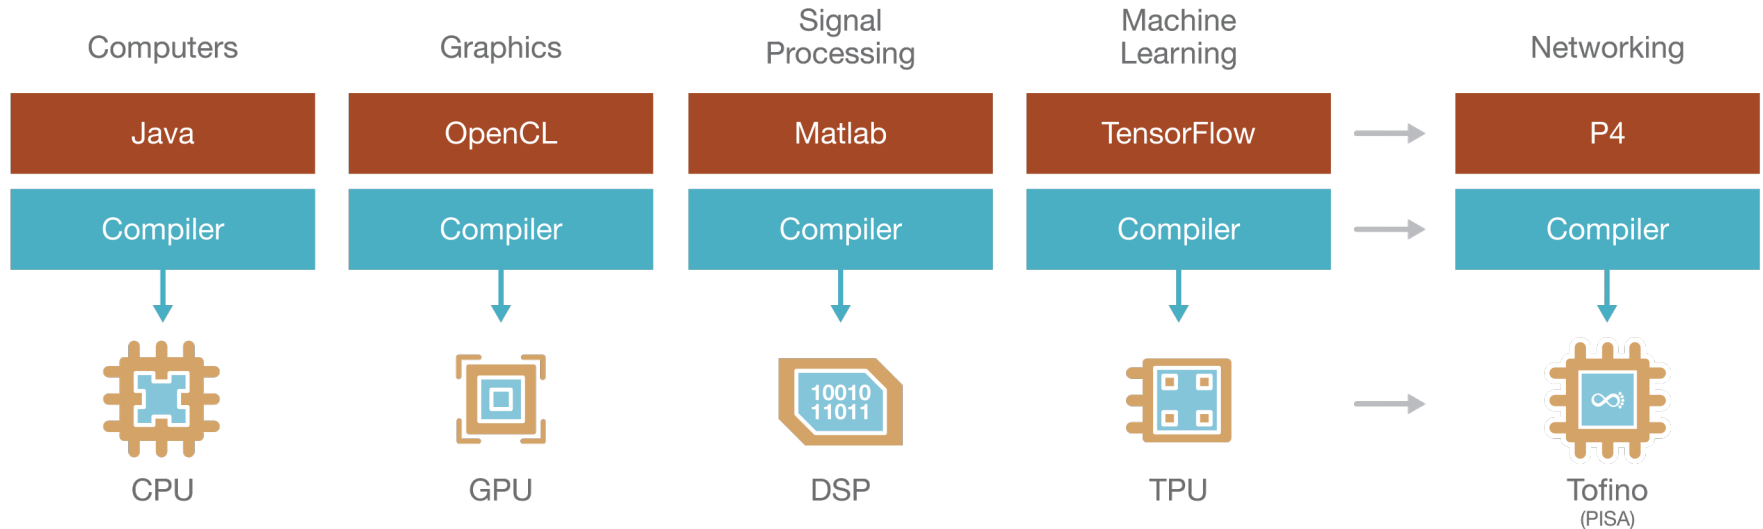
- **Open source**, evolving, domain-specific language
- Permissive Apache license, code on GitHub today

- **Membership is free**: contributions are welcome
- Independent, set up as a California nonprofit

# The P4 language - general information

*„Our goal is for P4<sub>16</sub> to enable the same kind of programmability for network data-planes as the CUDA language did for graphics cards”*

*Mihai Budiu, Chris Dodd, „The P4<sub>16</sub> Programming Language”*



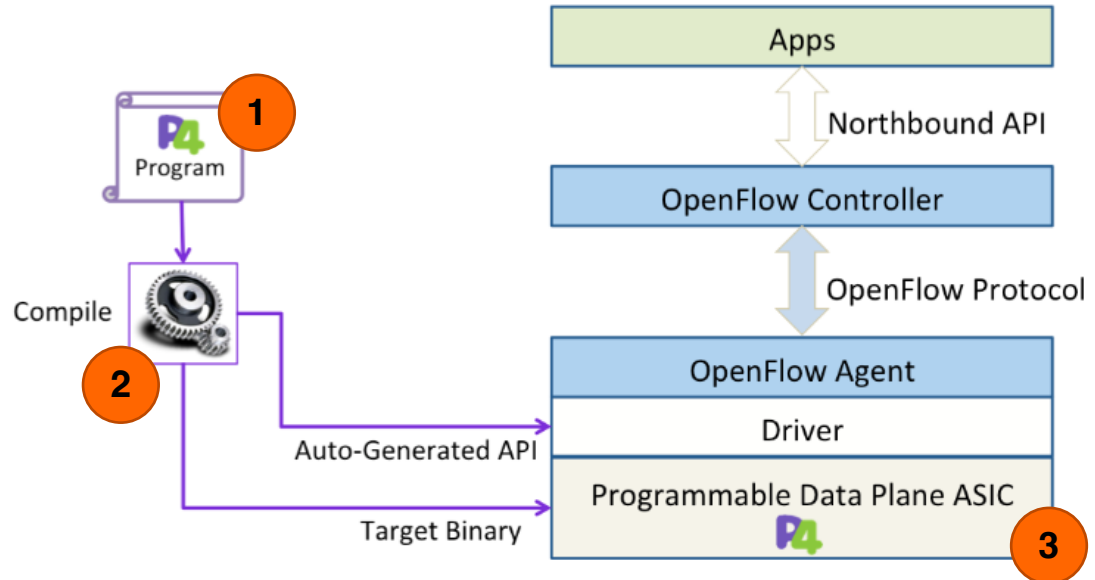
# The P4 technology

## 1. The P4 language – declarative Domain-Specific Language (DSL) to define data plane behavior

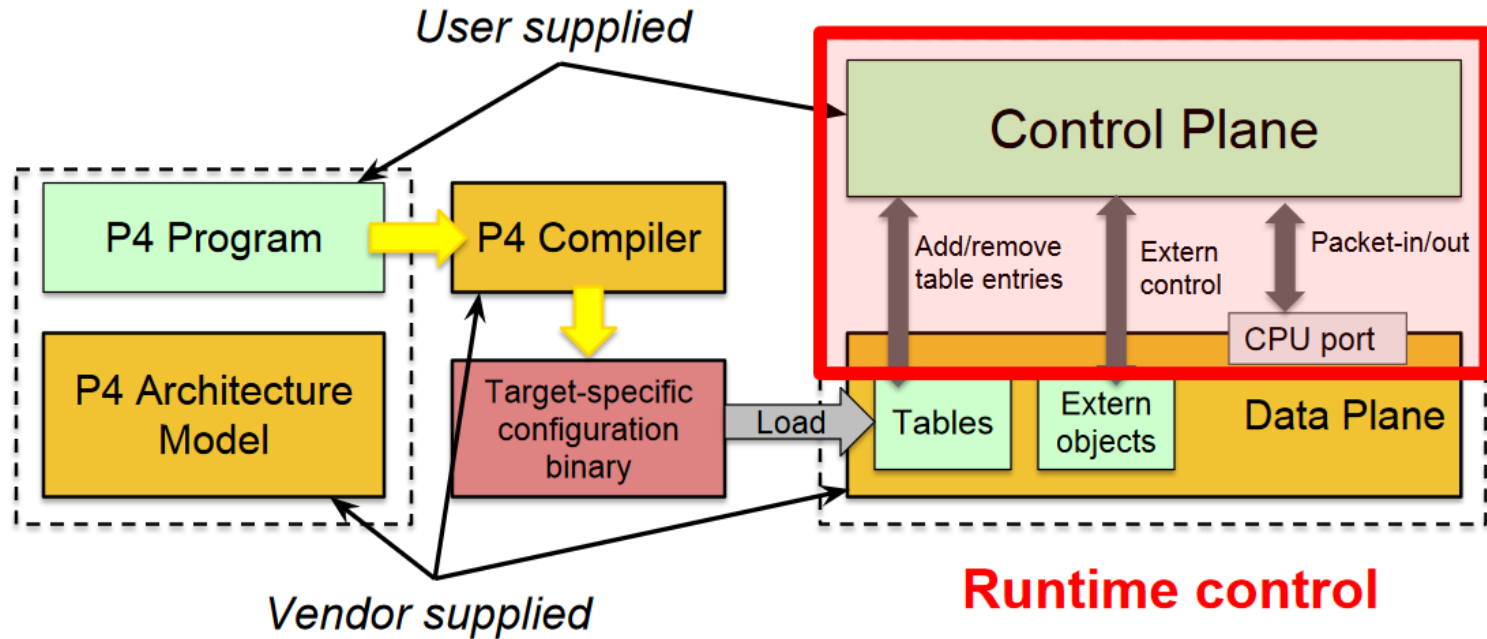
- Protocol-independent
- Target-independent
- Field Reconfigurable

## 2. The P4 Compiler

## 3. Abstract forwarding model for the network device

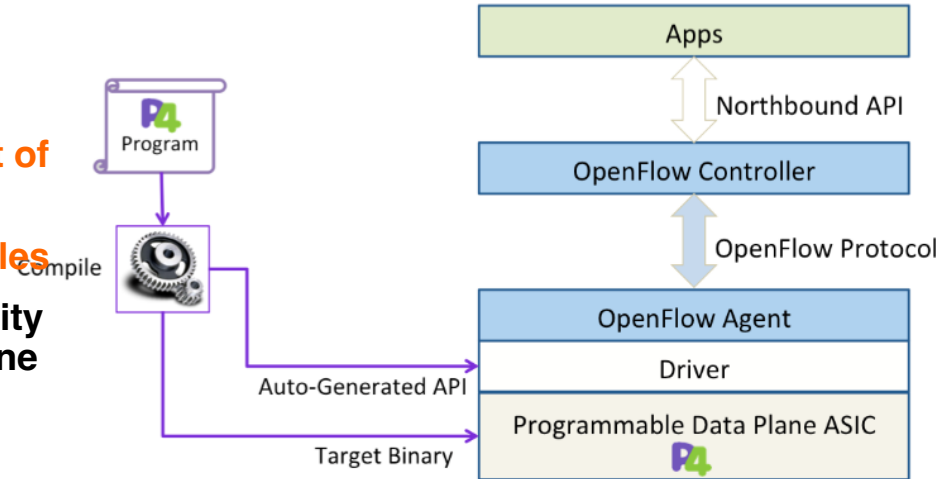


# Typical P4 workflow



# P4 vs. OpenFlow

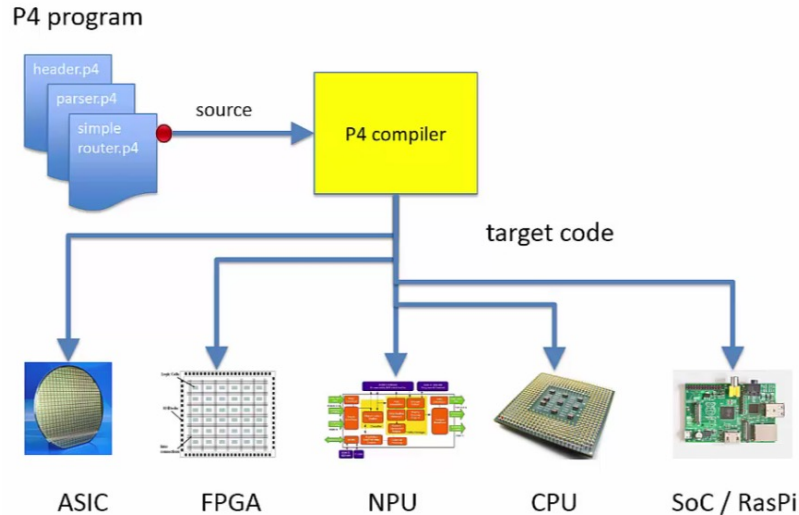
- What is the relationship between P4 and OpenFlow?
  - Both are invented to create open APIs for network devices
  - Serves different roles:
    - OpenFlow gives the way to populate a set of well-known tables
    - P4 gives the way to define forwarding tables
  - OpenFlow is more focused on programmability of control plane, while P4 focuses on data plane
  - OpenFlow assumes fixed instruction set for switch, while P4 allows to program that instruction set
- „While OpenFlow is designed for SDN networks in which we separate the control plane from the forwarding plane, P4 is designed to program the behavior of *any* switch or router”
- Question: Why are switch chips fixed-function, so far?



More on: <https://p4.org/p4/clarifying-the-differences-between-p4-and-openflow.html>

## P4: Target-independence

- **P4 Target – the platform, where the P4 program is executed on.**
  - Target's vendor must provide a corresponding compiler for its platform.
- **Target-independence – P4 language is independent from underlying platform (as Java)**
  - Compiler should translate a *target-independent* description (P4 program) into *target-dependent* code (used to configure the switch)



# P4: Protocol-independence

- P4 language can be used to define any network protocol
  - P4 has no native support for any protocol (such as Ethernet, IP, etc.)
  - The programmer must define all headers and appropriate parsers

```
typedef bit<48> macAddr_t;  
  
header ethernet_t {  
    macAddr_t dstAddr;  
    macAddr_t srcAddr;  
    bit<16> etherType;  
}
```

Ethernet

```
typedef bit<32> ip4Addr_t;  
  
header ipv4_t {  
    bit<4> version;  
    bit<4> ihl;  
    bit<8> diffserv;  
    bit<16> totalLen;  
    bit<16> identification;  
    bit<3> flags;  
    bit<13> fragOffset;  
    bit<8> ttl;  
    bit<8> protocol;  
    bit<16> hdrChecksum;  
    ip4Addr_t srcAddr;  
    ip4Addr_t dstAddr;  
}
```

IPv4

```
header_type icmp_t {  
    fields {  
        typeCode : 16;  
        hdrChecksum : 16;  
    }  
}
```

ICMP

```
header_type tcp_t {  
    fields {  
        srcPort : 16;  
        dstPort : 16;  
        seqNo : 32;  
        ackNo : 32;  
        dataOffset : 4;  
        res : 4;  
        flags : 8;  
        window : 16;  
        checksum : 16;  
        urgentPtr : 16;  
    }  
}
```


TCP

```
header_type udp_t {  
    fields {  
        srcPort : 16;  
        dstPort : 16;  
        length_ : 16;  
        checksum : 16;  
    }  
}
```

UDP

## P4: Reconfigurability

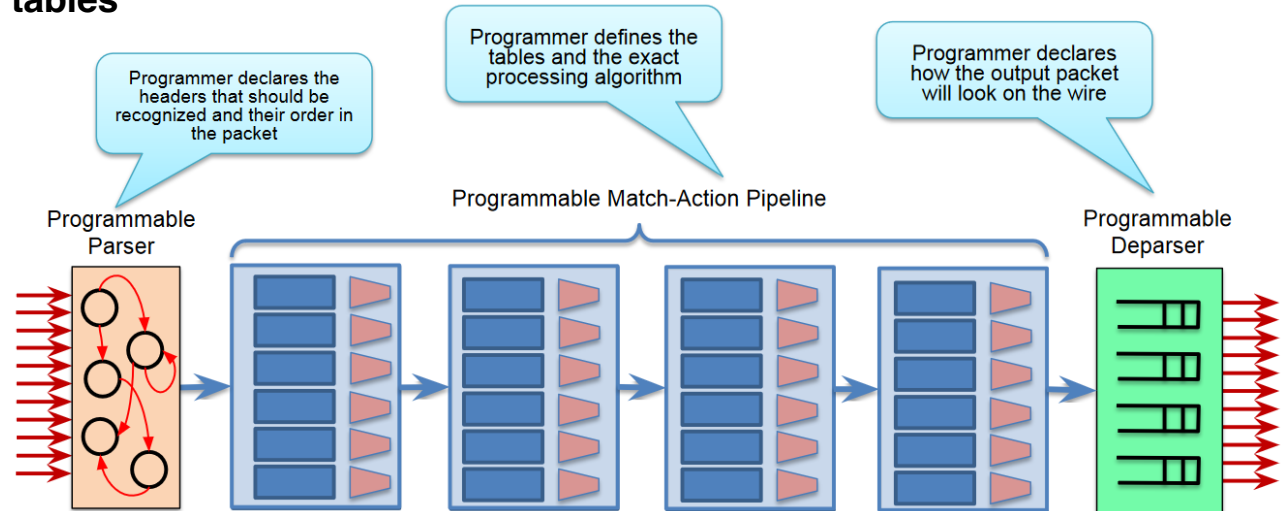
- It should be possible to re-define packet processing behavior on-the-fly.
- P4 has been designed to support reconfigurability
  - P4 provides high-level abstract model
  - P4 Target should be able to change the way of processing packets
- However, not all targets are re-configurable (e.g. fixed ASICs)





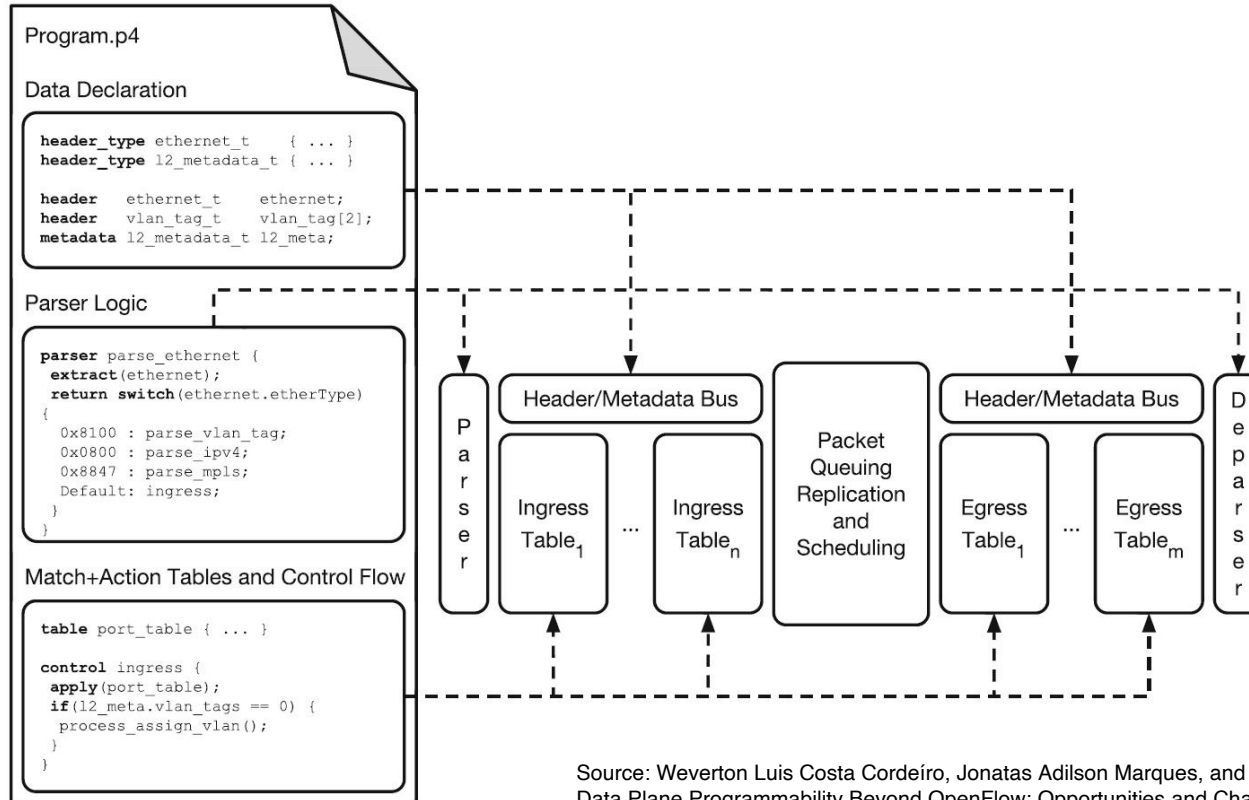
# P4 language – architecture model

- **Headers** – definitions of packet's headers
- **Parser** – finite state machine, tells how to process incoming bytes and extract headers
- **Deparser** – defines how the output packet will look like on the wire
- **Match-Action tables** – lookup keys + corresponding actions
- **Control flow** - sequence of tables



P4 Switch Architecture - Abstraction Model

# The P4 switch model



Source: Weverton Luis Costa Cordeiro, Jonatas Adilson Marques, and Luciano Paschoal Gaspar. 2017. Data Plane Programmability Beyond OpenFlow: Opportunities and Challenges for Network and Service Operations and Management. *J. Netw. Syst. Manage.* 25, 4 (October 2017)

# P4 Language – Types & Headers (1/2)

```
typedef bit<48> macAddr_t;  
typedef bit<32> ip4Addr_t;
```


```
header ethernet_t {  
    macAddr_t dstAddr;  
    macAddr_t srcAddr;  
    bit<16> etherType;  
}
```

```
header ipv4_t {  
    bit<4> version;  
    bit<4> ihl;  
    bit<8> diffserv;  
    bit<16> totalLen;  
    bit<16> identification;  
    bit<3> flags;  
    bit<13> fragOffset;  
    bit<8> ttl;  
    bit<8> protocol;  
    bit<16> hdrChecksum;  
    ip4Addr_t srcAddr;  
    ip4Addr_t dstAddr;  
}
```

- **Basic types:**
  - **bit<n>** - unsigned integer of size „n”
  - **int<n>** - signed integer of size n (<=2) // **Why?**
  - **varbit<n>** - variable-length field (e.g. for IPv4 options)
- **Headers:**
  - Defines the structure of protocol's header
  - Group of bit<n> fields
  - Defined based on specification e.g. RFCs
  - Can be valid or invalid
- **Typedef** – the construct to define a new type

## P4 – Protocol parser

- **Parser** – the construct that map packets into headers and metadata
  - Copy header's values from packet buffer to local memory
  - Finite state machine (FSM)
- **Every parser has three predefined states:**
  - *start*
  - *accept*
  - *reject*
- **If protocol is not supported by parser -> reject packet**



## P4 – Protocol parser example

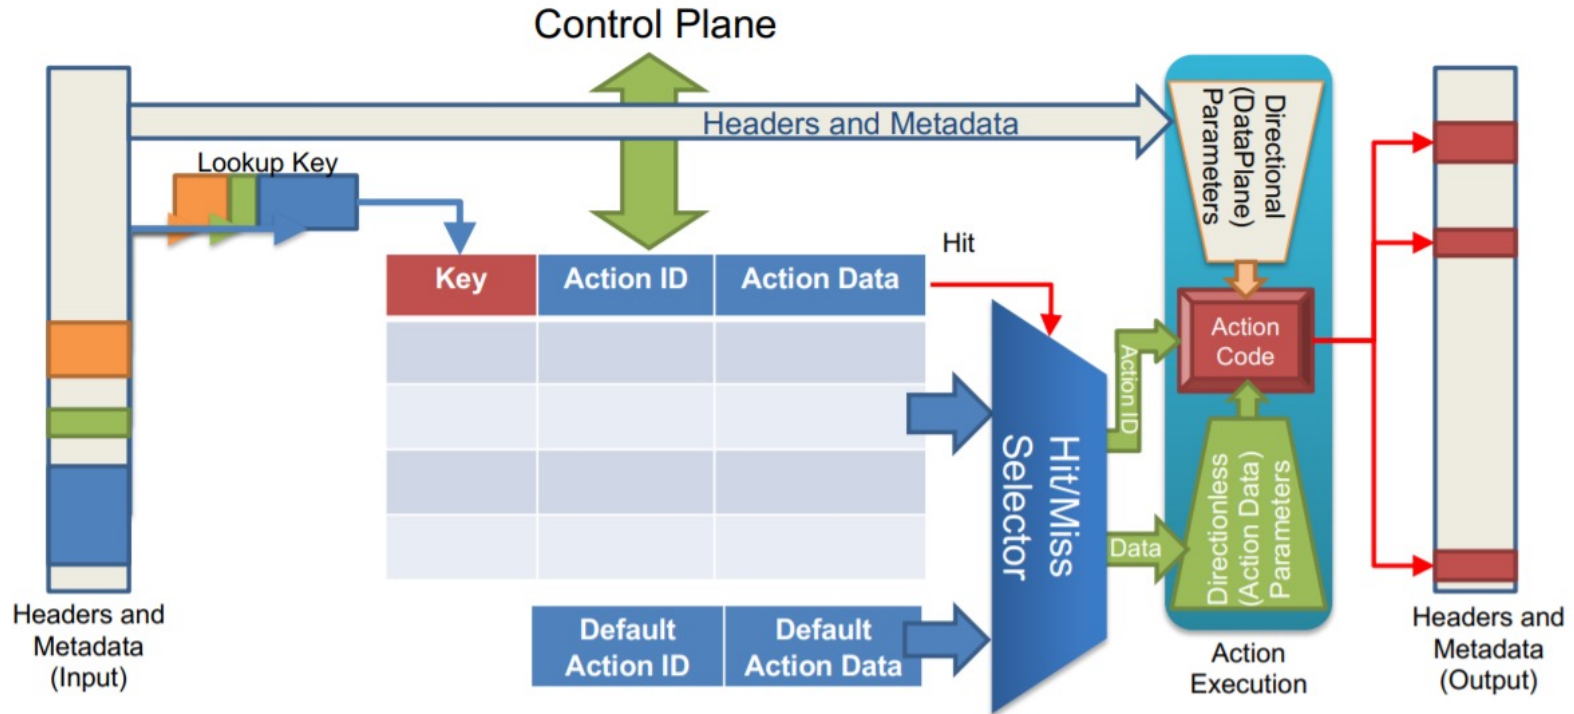
```
parser ParserImpl(packet_in packet,
                  out headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t std_meta) {

    state start {
        transition parse_ethernet;
    }


    state parse_ethernet {
        packet.extract(hdr.ethernet);
        transition select(hdr.ethernet.etherType) {
            0x0800: parse_ipv4;
            0x0806: parse_arp;
            0x86DD: parse_ipv6;
            0x8847: parse_mpls;
            default: reject;
        }
    }
}
```

- **Parser keywords:**
  - **select()** – used to branch in a parser
    - *similar to switch-case statement from C or Java*
  - **transition** – used to go to the next parser
- **Implementing parser requires expertise in network protocols:**
  - Next stage in parser is determined by the protocol header's field
  - For example, etherType is used to determine next protocol for Ethernet frames

## P4 - Match-Action tables (1/3)



## P4 - Match-Action tables (2/3)



Key	Action	Action Data
10.0.1.1/32	ipv4_forward	dstAddr=00:00:00:00:01:01 port=1
10.0.1.2/32	drop	
*^	NoAction	

- **Data plane (P4) program:**
  - Defines the format of the table:
    - Key fields
    - Actions + Action Data
  - Performs lookup
  - Executes the chosen action
- **Control plane (IP stack, routing protocols, SDN controller)**
  - Populates table entries with specific information based on:
    - Configuration (e.g. BGP routing policy)
    - Automatic discovery (e.g. MAC learning)
    - Protocol calculations (Dijkstra algorithm in OSPF)

## P4 - Match-Action tables (3/3)

```
/* Defined in core.p4 */
match_kind {
    exact,
    ternary,
    lpm
}
action NoAction() { }
action drop() {
    mark_to_drop();
}
action MyAction(macAddr_t dstAddr) { ... }

table ipv4_lpm {
    key = {
        hdr.ipv4.dstAddr: lpm;
    }
    actions = { ipv4_forward;
                drop;
                NoAction; }
    size = 1024;
    default_action = NoAction();
}
```

- **Table keywords:**
  - **key** – defines set of match fields
  - **actions** – defines a list of supported actions
  - **size** – defines a maximum number of entries
  - **default\_action** – defines an action to invoke if there is no match
- **Different match types:**
  - **exact**
  - **lpm** (Longest-Prefix Match)
  - **ternary**
- **Actions declared as simple functions in other languages (e.g. C, Java, etc.)**



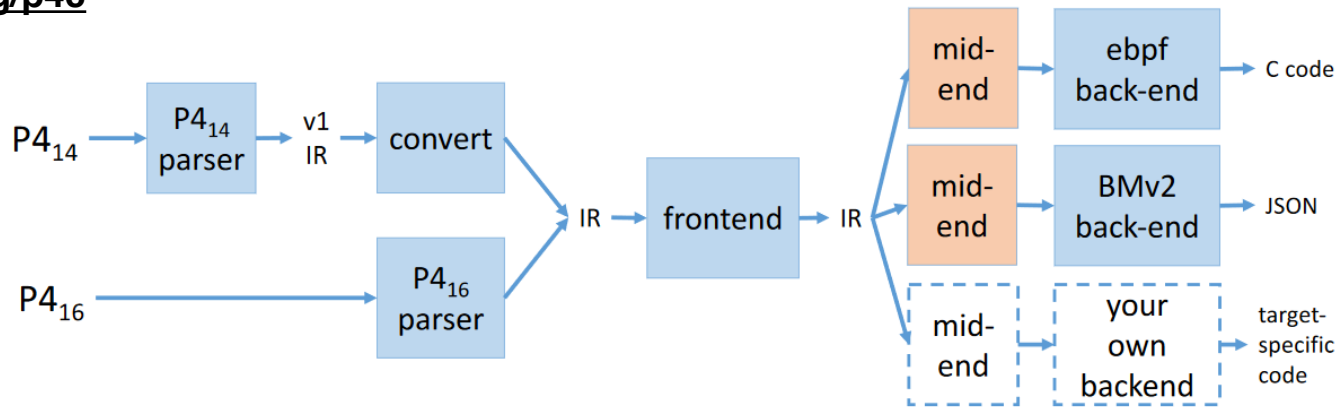
## P4 - Deparser

```
control deparser(packet_out packet,  
    in headers hdr) {  
  
    apply {  
        packet.emit(hdr.ethernet);  
        packet.emit(hdr.ipv4);  
    }  
}
```

- **Assembles the headers back into a well-formed packet**
- **emit() function serializes the header if it is valid**


# The P4 compiler

- **Compiler goals:**
  - Support current and future P4 versions
  - Open-source front-end
  - Support for multiple back-ends
    - Generate code for ASICs, NICs, FPGAs, software switches and other targets
- <https://github.com/p4lang/p4c>



# P4 targets – hardware platforms


- **NetFPGA**
  - Prototyping and evaluating P4 program on real hardware
  - <https://github.com/NetFPGA/P4-NetFPGA-public/wiki>
- **Barefoot Tofino**
  - Production-ready P4 switch by Barefoot
  - 12.8 Tb/s
- **SmartNICs, IPU, DPU**



Intel IPU E2000



Barefoot Tofino


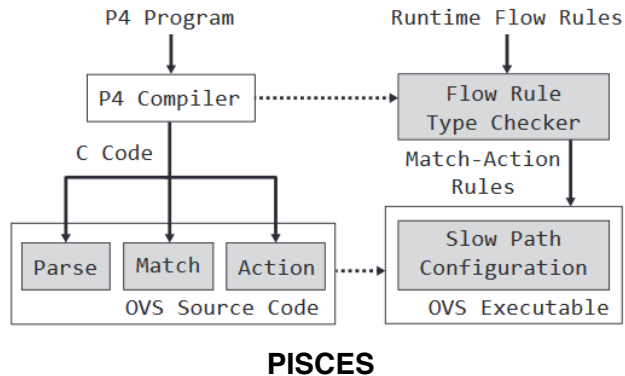


NetFPGA



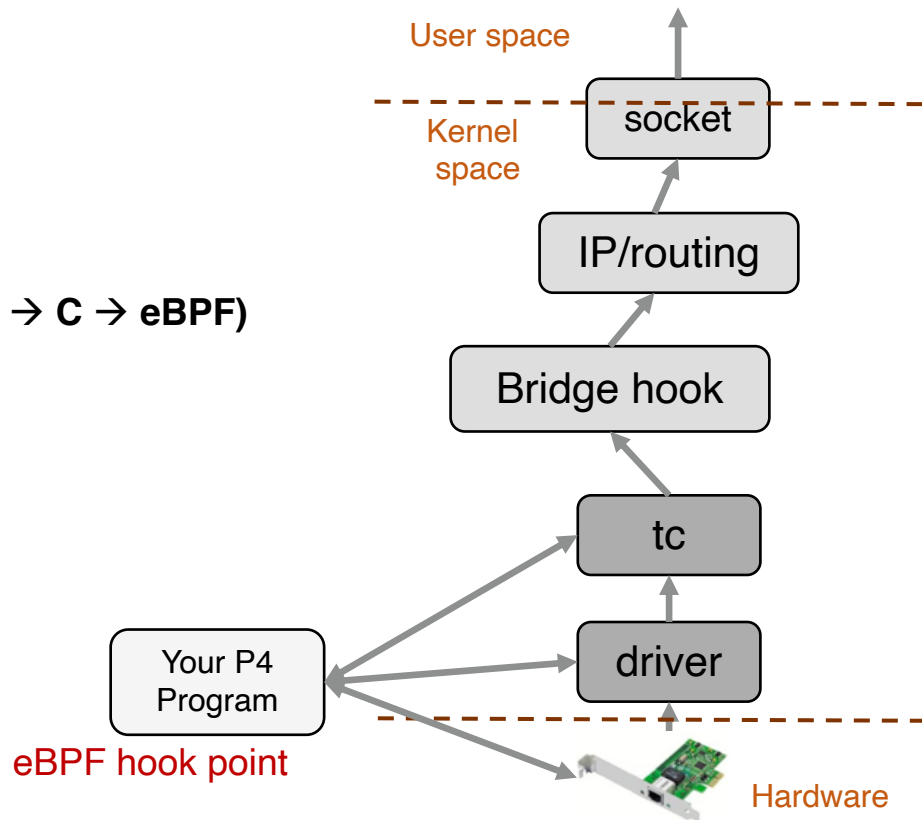

# P4 targets – software platforms

- **BMv2 (Behavioral Model v2)**
  - The P4 software reference switch
  - Up-to-date with P4 specification
- **PISCES**
  - P4-compatible Open vSwitch
  - P4-to-OVS compiler
- **P4-to-eBPF**



# P4-to-eBPF

- **eBPF (extended Berkeley Packet Filter)**
  - In-kernel Virtual Machine for packet filtering
  - Injecting code in the kernel at runtime
- **P4-to-eBPF**
  - Compiles P4 program to eBPF bytecode ( $P4 \rightarrow C \rightarrow eBPF$ )



Example of TC+eBPF

# Coding break: simple\_l3.p4



## simple\_l3.p4 – Basic structure (template)

```
#include <core.p4>
#include <v1model.p4>

/** CONSTANTS AND TYPES */

/** HEADER DEFINITIONS */

parser MyParser() { ... }

control MyIngressControl() { ... }

control MyEgressControl() { ... }

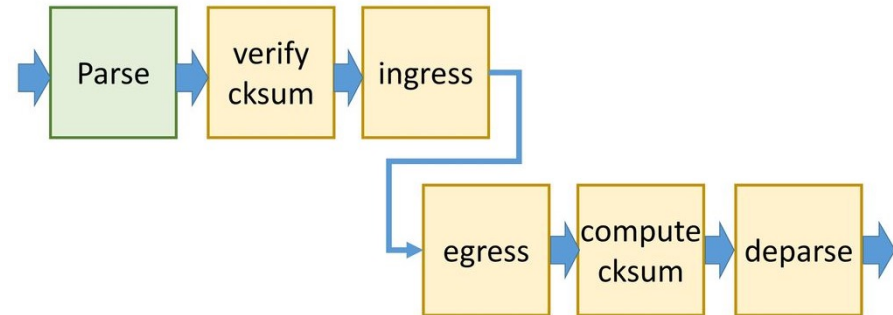
control MyDeparser() { ... }

control MyComputeChecksum() { ... }

control MyVerifyChecksum() { ... }

V1Switch(MyParser(), MyVerifyChecksum(),
         MyIngressControl(), MyEgressControl(),
         MyComputeChecksum(), MyDeparser()) main;
```

- We will use **v1model.p4**
- Write code for all P4-programmable components, allowed by the architecture model
  - *Your* Headers and Metadata
  - *Your* Parser
  - *Your* control blocks
  - *Your* logic to verify & compute checksum
  - *Your* Deparser
- Assemble everything in the **V1Switch** package



## simple\_l3.p4 – Constants, types and headers

```
#include <core.p4>
#include <v1model.p4>
```

```
/** CONSTANTS AND TYPES */
```

```
const bit<16> TYPE_IPV4 = 0x0800;
```

```
typedef bit<9> egressSpec_t;
typedef bit<48> macAddr_t;
typedef bit<32> ip4Addr_t;
```

```
/** HEADER DEFINITIONS */
```

```
header ethernet_t {
    macAddr_t dstAddr;
    macAddr_t srcAddr;
    bit<16> etherType;
}
```

```
header ipv4_t {
    bit<4> version;
    bit<4> ihl;
```

```
    bit<8> diffserv;
    bit<16> totalLen;
    bit<16> identification;
    bit<3> flags;
    bit<13> fragOffset;
    bit<8> ttl;
    bit<8> protocol;
    bit<16> hdrChecksum;
    ip4Addr_t srcAddr;
    ip4Addr_t dstAddr;
}
```

```
struct headers {
    ethernet_t ethernet;
    ipv4_t     ipv4;
}
```

```
struct routing_metadata_t { ip4Addr_t nhop_ipv4; }
```

```
struct metadata { routing_metadata_t routing; }
```



## simple\_l3.p4 – Parser & Deparser

```
parser RouterParser(packet_in packet,
                    out headers hdr,
                    inout metadata meta,
                    inout standard_metadata_t std_meta)
{
    state start {
        transition parse_ethernet;
    }
    state parse_ethernet {
        packet.extract(hdr.ethernet);
        transition select(hdr.ethernet.etherType) {
            TYPE_IPV4: parse_ipv4;
            default:    reject;
        }
    }
    state parse_ipv4 {
        packet.extract(hdr.ipv4);
        transition accept;
    }
}
```

```
control RouterDeparser(packet_out packet,
                       in headers hdr) {
    apply {
        packet.emit(hdr.ethernet);
        packet.emit(hdr.ipv4);
    }
}
```

- **Parser:**
  - Parse Ethernet -> IPv4
- **Deparser:**
  - deparse packets in reversed order

## simple\_l3.p4 – Ingress control block

```
action drop() { mark_to_drop() };
control ingress(...) {
    action ipv4_forward(ip4Addr_t nextHop,
                        egressSpec_t port) {
        meta.routing.nhop_ipv4 = nextHop;
        standard_metadata.egress_spec = port;
        hdr.ipv4.ttl = hdr.ipv4.ttl - 1;
    }

    table routing_table {
        key = { hdr.ipv4.dstAddr: lpm; }
        actions = { ipv4_forward;
                    drop;
                    NoAction; }
        default_action = NoAction();
    }

    apply {
        if (hdr.ipv4.isValid())
            routing_table.apply();
    }
}
```

- The role of **Ingress** control block is to implement L3 functionality:
  - Match packets based on IPv4 destination address (LPM)
  - Determine IP next hop (an egress port)
  - Decrement TTL
- Egress port must be determined in the Ingress pipeline
- **meta.routing.nhop\_ipv4** is used to store next hop IPv4 (to be used in the **Egress** pipeline)
- Specifics of v1model / BMv2:
  - **standard\_metadata.t.egress\_spec** stores the output port

## simple\_l3.p4 – Egress control block

```
control egress(...) {  
  action set_dmac(macAddr_t dstAddr) {  
    hdr.ethernet.dstAddr = dstAddr;  
  }  
  action set_smac(macAddr_t mac) {  
    hdr.ethernet.srcAddr = mac;  
  }  
  table switching_table {  
    key = { meta.routing.nhop_ipv4 : exact; }  
    actions = { set_dmac; }  
    default_action = NoAction();  
  }  
  table mac_rewriting_table {  
    key = { standard_metadata.egress_port: exact; }  
    actions = { set_smac; }  
    default_action = drop();  
  }  
  apply {  
    switching_table.apply();  
    mac_rewriting_table.apply();  
  }  
}
```

- The role of **Egress** control block is to implement L2 functionality
  - Set source MAC address equal to the egress port
  - Set destination MAC address equal to the MAC address of next-hop interface
- **meta.routing.nhop\_ipv4** is used to lookup the destination MAC address of the next hop

## simple\_l3.p4 – Verify & compute checksum

```
control MyComputeChecksum(inout headers hdr,
                          inout metadata meta) {

    apply {
        update_checksum(
            hdr.ipv4.isValid(),
            { hdr.ipv4.version,
              hdr.ipv4.ihl,
              hdr.ipv4.diffserv,
              hdr.ipv4.totalLen,
              hdr.ipv4.identification,
              hdr.ipv4.flags,
              hdr.ipv4.fragOffset,
              hdr.ipv4.ttl,
              hdr.ipv4.protocol,
              hdr.ipv4.srcAddr,
              hdr.ipv4.dstAddr },
            hdr.ipv4.hdrChecksum,
            HashAlgorithm.csum16);
    }
}
```

```
control MyVerifyChecksum(inout headers hdr,
                         inout metadata meta) {

    apply { }


}
```

- **Verify Checksum:**
  - Empty, don't verify checksum
- **Compute Checksum:**
  - Built-in function – `update_checksum()`
    - **1st arg:** condition to compute checksum
    - **2nd arg:** fields to compute checksum from
    - **3rd arg:** where to store checksum
    - **4th arg:** what algorithm to use (provided)


# Demo

- P4
- Mininet
- BMv2 (P4 software switch)

<https://github.com/P4-Research/p4-demos/tree/master/ip-routing>

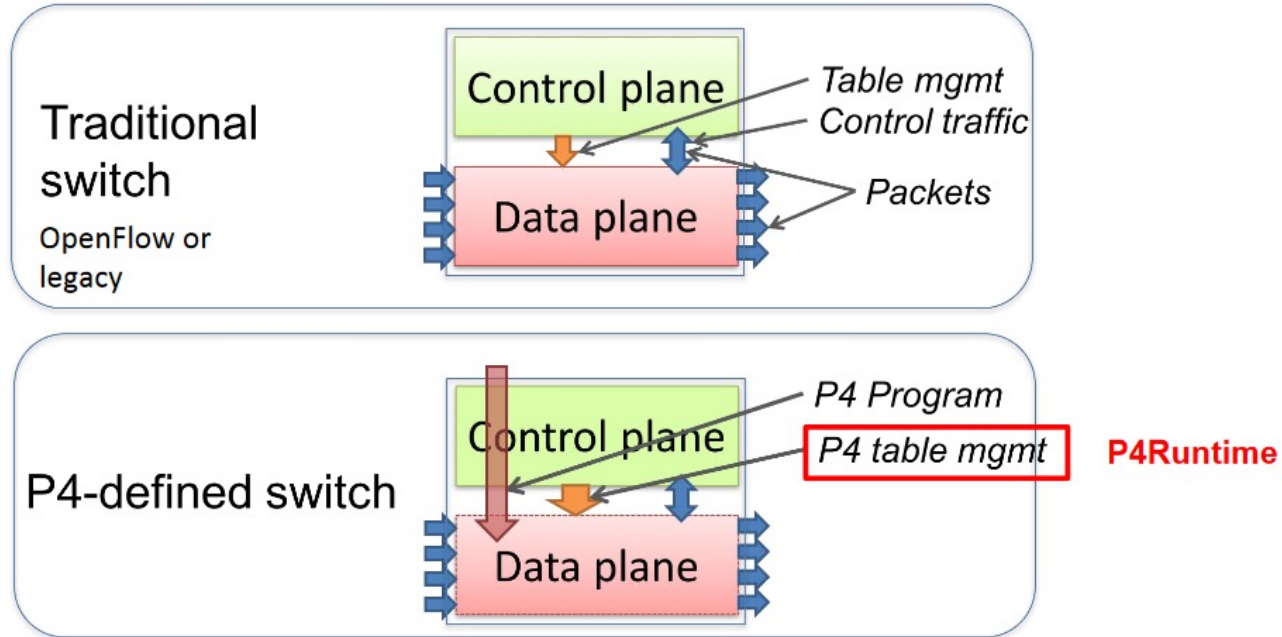


> sudo mn



**P4Runtime – a control plane for P4**


# Traditional/OpenFlow vs. P4 paradigm



# P4Runtime – how it is implemented?

- **SDN-like control plane for P4 switches**
  - Data plane and control plane decoupling
  - Logically centralized
  - Open API to data plane
- **Technologies:**
  - Google Protocol Buffers
  - gRPC

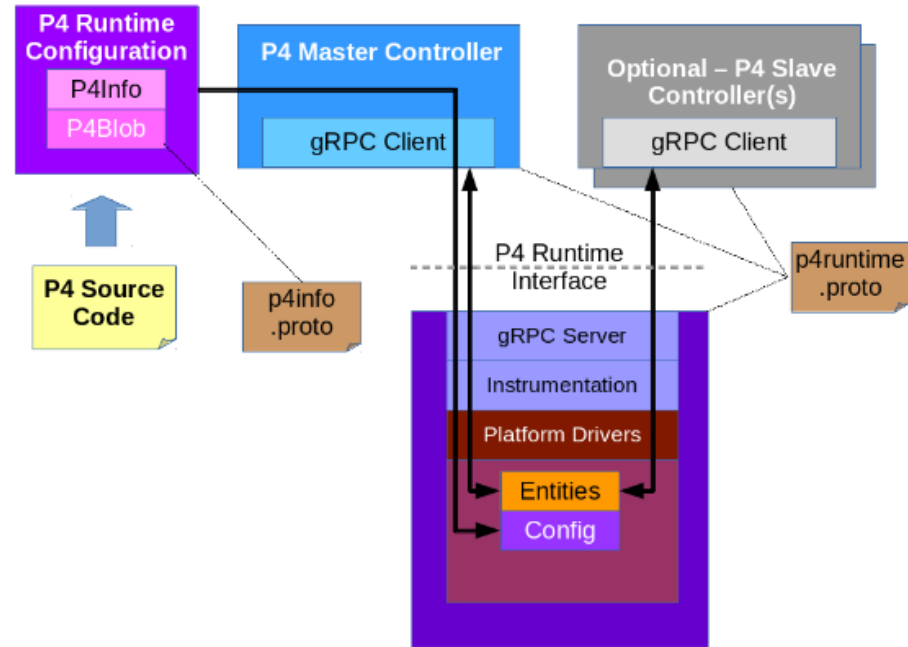
## P4 Runtime for remote Control Plane





# The P4Runtime protocol

- **The P4 compiler produces:**
  - ForwardingPipelineConfig
  - P4Info metadata – describes the structure of the P4 program (Tables, Actions, etc.)
- **The P4Runtime controller can:**
  - Set/Get ForwardingPipelineConfig
  - Write/Read Entities: TableEntry, Counter, etc.
- **P4Runtime's operations (RPCs):**
  - SetForwardingPipelineConfig()
  - GetForwardingPipelineConfigRequest()
  - Write()
  - Read()



# Support for P4Runtime in SDN controllers

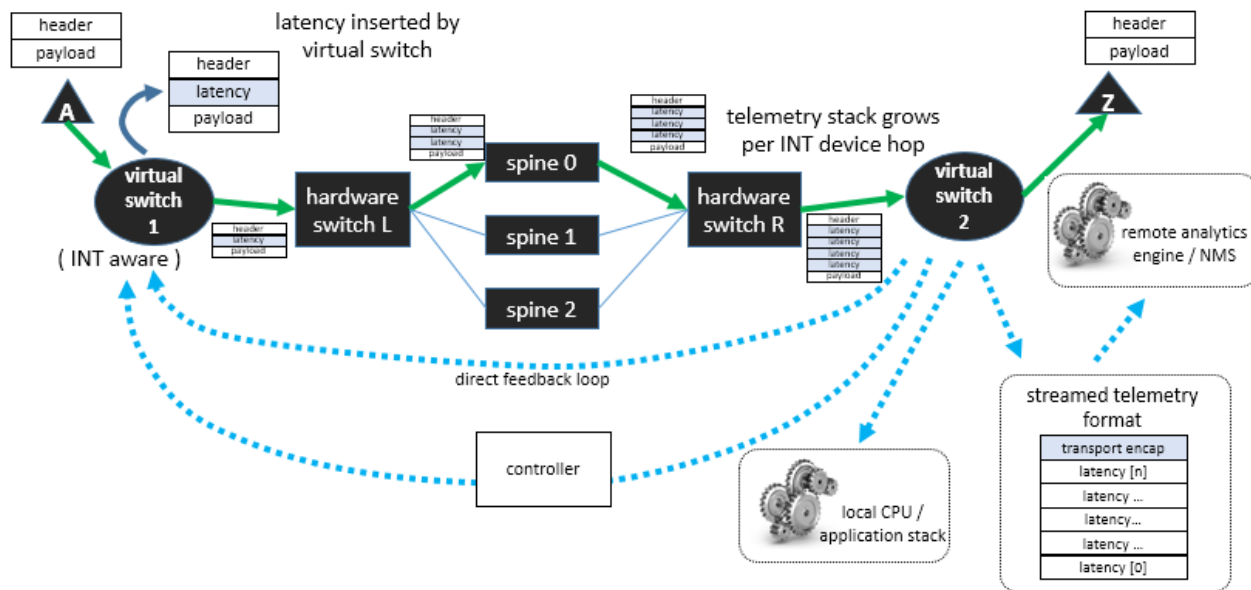
- **Open Network Operating System (ONOS)**
  - PoCs with **BMv2**, **Barefoot Tofino**, **Mellanox Spectrum** switch
  - <https://wiki.onosproject.org/display/ONOS/P4+brigade>
  - <https://wiki.onosproject.org/display/ONOS/P4Runtime+support+in+ONOS>
- **OpenDayLight**
  - [https://wiki.opendaylight.org/view/P4\\_Plugin:Main](https://wiki.opendaylight.org/view/P4_Plugin:Main) (led by ZTE)



# Use cases

# Use case study – In-band Network Telemetry (INT)

- The new monitoring tool
- Why P4 helps?



# Sample P4 applications

- **In-band Network Telemetry (INT)**
- **Network applications:**
  - **NAT, L3/L4 Load Balancing**
  - **Firewall**
  - **Segment Routing**
  - **Many others...**
- **VNF offloading**
  - **S-/P-GW (LTE)**
  - **Residential BNG (Broadband Network Gateway) with PPPoE termination**

# Research activities

- **P4 Architecture Working Groups**
- **Programmable QoS scheduling & traffic management**
- **New use cases for P4**
  - **Telemetry & Monitoring**
  - **VNF's implementations (NFV Acceleration)**
  - **New architectures for Data Center**
  - **Information-Centric Networking**
  - **Security**
  - **5G (Network Slicing)**
- **More on:**
  - **<https://open-nfp.org/projects/>**

# Summary

# Summary

- **Network innovation (evolution) is limited due to OpenFlow's limitations**
  - **Fixed data plane**
  - **New data plane protocol requires changes to specification & implementation**
- **Data plane programmability becomes a next step in the SDN evolution**
  - **Allows to re-configure data plane functions on-the-fly**
- **P4 is an enabler for data plane programmability. It provides:**
  - **Architecture for programming data plane**
  - **High-level language to describe data plane behavior**
  - **Abstract model for programmable switch**
- **P4 is promising, but still under development!**





**Questions?**